

1.

a) Attempted transformations

All the 'parallel' for is acted on j loop. We can't put 'parallel for' on i or k loop, because the dependence. All the transformations are based on this parallel.

I tried permutation ikj, and that will help a lot. So the further transformations are based on ikj. Then I tried tiled (IT-256), tiled(IT-8),tiled(KT-4),tiled(KT-8),tiled(KT-256), tiled(IT-256,JT-256), tiled(IT-256,JT-256,KT-256),tiled(IT-128,KT-128), tiled(IT-128,KT-256),tiled(IT-256,KT-256),tiled(IT-512,KT-512), tiled unroll(IT-256,KT-256,unroll-2j), tiled(JT-256,KT-512),tiled(JT-256,KT-256),unroll(2j), unroll(2i).

b) How performance varied with the different transformations.

for convenience, I just show the performance for number of threads=1,2,4,8,10.

For gcc:

threads	1	2	4	8	10
base	0.2	0.2	0.2	0.2	0.2
Just parallel	0.2	0.3	0.5	0.7	0.7
ikj	1.9	2.3	2.5	2.3	2.2
tiled(IT-8)	1.9	2.3	2.3	2.4	2.3
tiled(IT-256)	1.9	2.3	2.6	2.3	2.8
tiled(KT-4)	2	2.5	2.8	2.9	2.6
tiled(KT-8)	2.1	2.6	3.2	3.3	3.1
tiled(KT-256)	2.2	2.7	3.2	3.1	3.3
tiled(IT-256,JT-256)	0.9	0.6	0.6	0.5	0.5
tiled(IJK-256)	1.1	0.7	0.6	0.5	0.5
tiled(IT-128,KT-128)	2.2	2.8	3.5	3.5	3.3
tiled(IT-128,KT-256)	2.2	2.8	3.5	3.5	3.1
tiled(IT-256,KT-256)	2.2	2.8	3.5	3.6	3.1
tiled(IT-512,KT-512)	2.2	2.8	3.5	3.7	3.4
tiledUnrol(IK-256,2j)	1.4	2	2.8	3.5	3.4
tiled(JT-256,KT-512)	1.1	0.7	0.8	0.6	0.5
tiled(JT-256,KT-256)	1.1	0.7	0.8	0.5	0.5
Unroll-2i	1.7	2.7	4.3	4	3.1
Unroll-2j	1.3	1.8	2.3	2.6	2.1

From the table above, we can conclude that unroll i is a good idea to optimize. And tile I,K is fine but not J, because we do "parallel for" for j loop.

For icc:

threads	1	2	4	8	10
base	0.2	0.2	0.2	0.2	0.2
Just parallel	0.2	0.3	0.6	0.7	0.7
ikj	2.7	2.3	2.3	2.1	1.9
tiled(IT-8)	2	2.7	2.6	2.5	2.1
tiled(IT-256)	2.7	2.3	2.4	2	1.8
tiled(KT-4)	3.1	4.3	3	2.3	2
tiled(KT-8)	3.4	3.7	3.1	2.8	2.5
tiled(KT-256)	3.6	4.4	3.5	3.3	2.8
tiled(IT-256,KT-256)	0.6	0.5	0.5	0.3	0.3
tiled(IJK-256)	3	0.8	0.5	0.5	0.4
tiled(IT-128,KT-128)	3.7	3.3	3.4	3.4	3
tiled(IT-128,KT-256)	3.7	3.4	3.5	2.8	2.1
tiled(IT-256,KT-256)	3.6	3.5	3.7	2.4	2.3
tiled(IT-512,KT-512)	2.7	3.4	3.6	3.3	2.8
tiledUnrol(IK-256,2j)	2.3	2.6	3	3.1	2.9
tiled(JT-256,KT-512)	3	0.7	0.5	0.4	0.4
tiled(JT-256,KT-256)	2.9	0.7	0.5	0.4	0.3
Unroll-2i	2.6	3.9	4.7	3.7	3.3
Unroll-2j	1.7	2.1	2.2	1.9	1.8

From the table above, we can conclude that tile k and tile i,k will be good. Also unrolling j for more threads is better.

c) the performance for the best version for each compiler with each of 1-12 threads

Threads	gcc	icc
1	2.2	3.7
2	2.8	4.4
3	3.1	4.3
4	4.3	4.7
5	4.0	4.7
6	4.1	4.5
7	3.8	4.1
8	4.0	3.7
9	3.4	3.5
10	3.4	3.3
11	2.8	3.0
12	2.4	2.9

d) the code for my best version
for both gcc and icc:

```
#pragma omp parallel private(i,k)
// Template version is intentionally made sequential
// Make suitable changes to create parallel version
{
    if (omp_get_thread_num()==0 && omp_get_num_threads()
!= nt)
        printf("Warning: Actual #threads %d differs from
requested number %d\n",omp_get_num_threads(),nt);

    //
    // Test version of code; initially just contains a
copy of base code
    // To be modified by you to improve performance
    //
    // #pragma omp master
    for (i=0; i<N; i+=2)
    {
#pragma omp single
        for(j=0; j<N; j++)
            B[i][j] += A[i][i] * B[i][j];
        for(k=i+1; k<N; k++)
#pragma omp for
            for(j=0; j<N; j++)
            {
                B[i][j] += A[i][k] * B[k][j];
                B[i+1][j] += A[i+1][k] * B[k][j];
            }
    }
} // End of parallel region
```

e) attachment

I put the attachment in the back, which contains all the results for all the number of threads.

2.

a) Attempted transformations

First of all, I did the permutation `klij` without the parallel. That was pretty faster than the old one. Then did the parallel. The parallel is that put “parallel for” on the outer loop of `k`. The further transformations are based on permutation and parallel. The attempted transformations are permutation `klij`, parallel for `klij`, `tile(i/2,j/2)`, `tile(i/4,j/4)`, `thileUnroll(i/2,j/2,2i)`, `unroll(2i)`, `unroll(4i)`, `unroll(2j)`, `unroll(4j)`.

b) How performance varied with the different transformations.

for convenience, I just show the performance for number of threads=1,2,4,8,12.

For gcc:

threads	1	2	4	8	12
base	0.2	0.2	0.2	0.2	0.2
<code>klij</code>	2.3	2.3	2.3	2.3	2.3
<code>klij-parallel</code>	2.3	3	4.4	5.9	5.5
<code>tile(i-2,j-2)</code>	1.7	2.1	3.5	5.5	6.3
<code>tile(i-4,j-4)</code>	1.6	2	2.6	4.6	5.8
<code>tileUnroll(i-2,j-2,2i)</code>	1.7	2.1	3.5	4.6	6.2
<code>unroll(2i)</code>	2.4	3.1	3.6	6.8	7.9
<code>unroll(4i)</code>	1.8	2.3	3.9	5.9	7
<code>unroll(2j)</code>	2.5	3.3	3.6	5.5	6.5
<code>unroll(4j)</code>	2.6	3.3	3.6	5.6	6.6

We got the best results by doing `unroll(2i)`.

For icc.

threads	1	2	4	8	12
base	0.3	0.3	0.3	0.3	0.3
<code>klij</code>	2.3	2.9	2.9	2.9	2.8
<code>klij-parallel</code>	2.4	3.4	5.3	6.7	6.1
<code>tile(i-2,j-2)</code>	1.9	2.9	4.5	6.8	7.5
<code>tile(i-4,j-4)</code>	1.6	2.1	3.3	5.8	6.5
<code>tileUnroll(i-2,j-2,2i)</code>	1.5	1.9	3.1	5.7	6.1
<code>unroll(2i)</code>	1.8	2.3	3.6	6.2	6.5
<code>unroll(4i)</code>	2.1	2.8	4.4	6.8	7.6
<code>unroll(2j)</code>	1.8	2.5	3.6	5.8	6.4
<code>unroll(4j)</code>	2.2	3.4	5	7.9	8.2

We got the best results by doing `unroll(4j)`.

c) the performance for the best version for each compiler with each of 1-12 threads

Threads	gcc	icc
1	2.6	2.4
2	3.3	3.4
3	3.8	4.3
4	4.4	5.3
5	4.5	6.1
6	5.8	6.6
7	6.3	7.6
8	6.8	7.9
9	6.8	7.8
10	7.6	7.8
11	7.9	8.0
12	7.9	8.2

d) the code for my best version.

For gcc:

```
#pragma omp parallel private(i,j,k,l)
// Template version is intentionally made sequential
// Make suitable changes to create parallel version
{
    if (omp_get_thread_num()==0 && omp_get_num_threads()
!= nt)
        printf("Warning: Actual #threads %d differs from
requested number %d\n",omp_get_num_threads(),nt);

    //
    // Test version of code; initially just contains a
copy of base code
    // To be modified by you to improve performance
    //
    // #pragma omp master
#pragma omp for
    for(k=0; k<N; k++)
        for(l=k+1; l<N; l++)
            for (i=0; i<N; i+=2)
                for (j=0; j<N; j++){
                    C[k][l] +=
0.5*(A[l][i][j]*B[k][i][j]+A[k][i][j]*B[l][i][j]);
                    C[k][l] +=
0.5*(A[l][i+1][j]*B[k][i+1][j]+A[k][i+1][j]*B[l][i+1][j]);
                }
    } // End of parallel region
```

for icc:

```
#pragma omp parallel private(i,j,k,l)
// Template version is intentionally made sequential
// Make suitable changes to create parallel version
{
    if (omp_get_thread_num()==0 && omp_get_num_threads()
!= nt)
        printf("Warning: Actual #threads %d differs from
requested number %d\n",omp_get_num_threads(),nt);

    //
    // Test version of code; initially just contains a
copy of base code
    // To be modified by you to improve performance
    //
    // #pragma omp master
#pragma omp for
    for(k=0;k<N;k++)
        for(l=k+1;l<N;l++)
            for (i=0; i<N; i++)
                for (j=0; j<N; j+=4){
                    C[k][l] +=
0.5*(A[l][i][j]*B[k][i][j]+A[k][i][j]*B[l][i][j]);
                    C[k][l] +=
0.5*(A[l][i][j+1]*B[k][i][j+1]+A[k][i][j+1]*B[l][i][j+1]);
                    C[k][l] +=
0.5*(A[l][i][j+2]*B[k][i][j+2]+A[k][i][j+2]*B[l][i][j+2]);
                    C[k][l] +=
0.5*(A[l][i][j+3]*B[k][i][j+3]+A[k][i][j+3]*B[l][i][j+3]);
                }
    } // End of parallel region
```

e) attachment

I put the attachment in the back, which contains all the results for all the number of threads.